

NTT DATA Sitecore CDN Connector

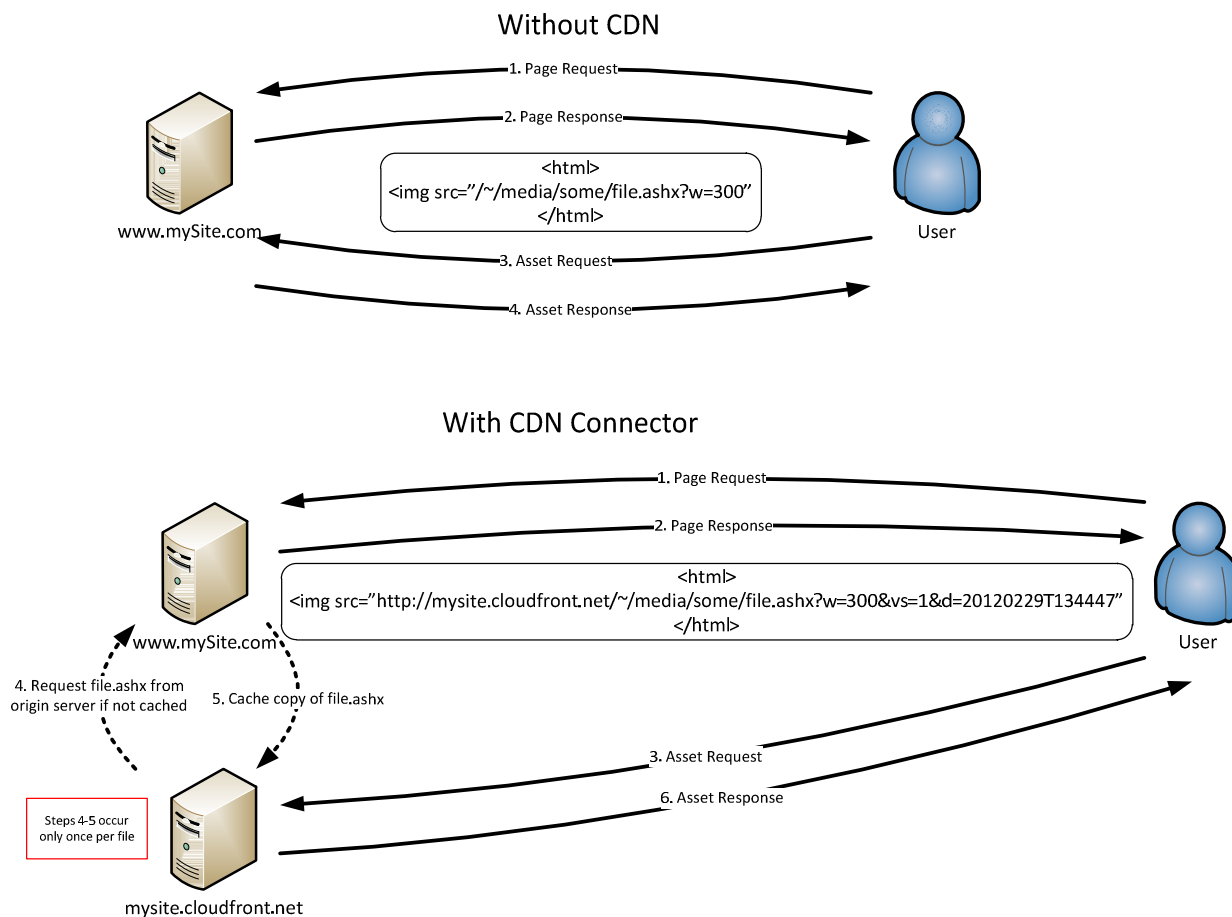
Overview

The CDN Connector for Sitecore allows developers to route all media requests (dynamic and static) through a proxy CDN. It is designed to be plug-n-play requiring no additional development effort other than configuration.

Media Rerouting

Media URLs that would traditionally be handled by the web server would be rewritten to point to a geographically appropriate edge server which will cache media items from the origin server as they are requested. This takes a large amount of load off the web (origin) server.

An example using AWS CloudFront would be:



Setup

The package simply needs to be installed and configured for the desired CDN's distribution server.

If working with AWS CloudFront, you can follow the below steps for setting up a distribution.

Setting up the distribution in AWS

1. Sign into the AWS console, then the CloudFront tab.
2. Click "Create Distribution"
3. Choose the Download delivery method
4. Provide the hostname of your public facing web server as the Origin Domain Name and accept the defaults

The screenshot shows the 'Create Origin' wizard in the AWS CloudFront console. The wizard has four steps: CHOOSE AN ORIGIN, CREATE DEFAULT BEHAVIOR, DISTRIBUTION DETAILS, and REVIEW. The first step, CHOOSE AN ORIGIN, is currently active. Below the step indicator, there is a description: 'An origin is the Amazon S3 bucket or the web server where you store the original versions of your web content. After you create your distribution, you can add more origins in the Distribution Details pane. [Learn More](#)'. The form contains the following fields: 'Origin Domain Name' with the value 'www.site.com', 'Origin ID' with the value 'Custom-www.site.com', 'Origin Protocol Policy' with radio buttons for 'HTTP Only' (selected) and 'Match Viewer', 'HTTP Port' with the value '80', and 'HTTPS Port' with the value '443'. At the bottom, there are '< Back' and 'Continue >' buttons.

5. Choose "Yes" for Forward Query Strings. The rest of the default values are fine

The screenshot shows the 'Create Cache Behavior' wizard in the AWS CloudFront console. The wizard has four steps: CHOOSE AN ORIGIN, CREATE DEFAULT BEHAVIOR, DISTRIBUTION DETAILS, and REVIEW. The second step, CREATE DEFAULT BEHAVIOR, is currently active. Below the step indicator, there is a description: 'A cache behavior determines how CloudFront communicates with your origin. After you create your distribution, you can add more cache behaviors in the Distribution Details pane. [Learn More](#)'. The form contains the following fields: 'Path Pattern' with the value 'Default (*)', 'Origin' with the value 'Custom-www.site.com', 'Viewer Protocol Policy' with radio buttons for 'HTTP and HTTPS' (selected) and 'HTTPS Only', 'Object Caching' with radio buttons for 'Use Origin Cache Headers' (selected) and 'Customize', 'Minimum TTL' with the value '0' and a help icon, and 'Forward Query Strings' with radio buttons for 'Yes' (selected) and 'No (Improves Caching)'. At the bottom, there are '< Back' and 'Continue >' buttons.

6. Accept the default values for the Distribution Details. CNAMEs are optional domain names you can set up from your DNS to point to the CloudFront domain name you'll receive later. (Example: "media.site.com")

The screenshot shows the 'Create Distribution' wizard in the AWS CloudFront console. The wizard has four steps: CHOOSE AN ORIGIN, CREATE DEFAULT BEHAVIOR, DISTRIBUTION DETAILS, and REVIEW. The third step, DISTRIBUTION DETAILS, is currently active. Below the step indicator, there is a description: 'Additional configuration is required. [?](#)'. The form contains the following fields: 'Alternate Domain Names (CNAMEs)' with the value 'media.site.com', 'Default Root Object' (empty), 'Logging' with radio buttons for 'On' and 'Off' (selected), 'Bucket for Logs' (empty), 'Log Prefix' (empty), 'Comments' (empty), and 'Distribution State' with radio buttons for 'Enabled' (selected) and 'Disabled'. At the bottom, there are '< Back' and 'Continue >' buttons.

7. Click through the rest of the wizard and you'll be returned to the distribution list

- Click on your new distribution and copy the "Domain Name". This is the publicly provided hostname if you are not using your own custom CNAMEs. If you are using custom CNAMEs use this "Domain Name" as the CNAME's target.
- You will have to wait for the distributions status to change from "In Progress" to "Enabled" before you can test this.

Installing on the sitecore application

Ensure the following files are installed

- /bin/NTTData.SitecoreCDN.dll
- /App_config/Include/SitecoreCDN.config

In your web.config, do the following:

- Under /system.webserver/handlers/
ADD

```
<add verb="*" path="aws_minify.ashx" type="NTTData.SitecoreCDN.Handlers.MinifyHandler, NTTData.SitecoreCDN" name="SitecoreCDN.Minify" />
```

 BEFORE

```
<add verb="*" path="sitecore_handlers.ashx" type="Sitecore.Web.CustomHandlerFactory, Sitecore.Kernel" name="Sitecore.GenericHandler"/>
```
- Under /system.web/httpHandlers/
ADD

```
<add path="aws_minify.ashx" verb="*" type="NTTData.SitecoreCDN.Handlers.MinifyHandler, NTTData.SitecoreCDN" />
```

 BEFORE

```
<add verb="*" path="sitecore_handlers.ashx" type="Sitecore.Web.CustomHandlerFactory, Sitecore.Kernel"/>
```
- In /App_Config/Includes/SitecoreCDN.config set the hostname of your CDN per site.

```
<sites>
  <site name="website">
    <patch:attribute name="cdnHostName">yourcdnhostname</patch:attribute>
  </site>
</sites>
```

Html Rewriting

The module rewrites the resulting html at the last moment by overriding the Response.Filter stream.

A pipeline processor `NTTData.SitecoreCDN.Pipelines.CDNAttachFilter` is located under

```
<httpRequestProcessed>
...
<processor type="NTTData.SitecoreCDN.Pipelines.CDNAttachFilter, NTTData.SitecoreCDN"/>
...
</httpRequestProcessed>
```

This pipeline decides whether to attach the Response.Filter based on a few criteria

- Sitecore.Context.Item != null (Item was resolved)
 - OR args.Url.FilePathWithQueryString matches an entry in <processRequets> config
- Sitecore.Context.Site != null (Context Site resolved)

3. Sitecore.Context.PageMode.IsNormal (Site is not being edited)
4. Sitecore.Context.Site has a “**cdnHostName**” attribute set (hostnames are per site)
5. Querystring url parameter “?cdn=true|false” (for overriding default behavior)

The rewriting filter then takes all the generated html from the page and replaces urls on these elements ``, `<script src="[url]" ></script>`, and `<link href="[url]" />` with transformed urls before returning the page to the client.

URL Rewriting

URLs are rewritten with 3 purposes:

1. Attaching the CDN hostname to make the url fully qualified and pointing to the proxy CDN
2. Attach version information and a datetime stamp to the query string to **guarantee uniqueness** so the CDN doesn't over-cache a file that's been modified.

Dynamic Media Urls

Dynamic urls are rewritten from

`/~/media/path/to/file.ashx?w=300&as=1` to

`http://cdnhostname/~/media/path/to/file.ashx?w=300&as=1&vs=3&d=20120101T120000`

The `&vs=3` is this media's sitecore Item version while `&d=20120101T120000` is the last modified time of the item for unversioned items. A CDN will treat these rewritten parameters as a unique filename when caching.

Static media Urls

Static urls are rewritten from

`/path/to/file.ext` to

`http://cdnhostname/path/to/file.ext?d=20120101T120000`

The `d=20120101T120000` is the “Last Modified” date of the file on the file system. A CDN will treat these rewritten parameters as a unique filename when caching.

Excluded Urls

Certain URLS will be excluded from being rewritten, instead keeping their local url.

1. Any external urls
2. Urls matching any `/sitecore/cdn/excludeUrls/regex` in `SitecoreCDN.config`
3. Media Items that are not publicly accessible to `[domain]\Anonymous`
4. Media Items that use Analytics Tracking

Consuming Rewritten URLs

The rewritten URL requests will go to the closest edge server. If the server hasn't previously cached this file, it will go to the origin server to get the first copy. Sitecore must be able to recognize these rewritten urls and recognize the embedded filepath and querystring parameters.

A `PreProcessRequest` pipeline step `NTTData.SitecoreCDN.Pipelines.CDNInterceptPipeline` is located under `<preprocessRequest>`

```
...
<processor type=" NTTData.SitecoreCDN.Pipelines.CDNInterceptPipeline, NTTData.SitecoreCDN"/>
...
```

```
</preprocessRequest>
```

The preprocessor will assess whether the item is a minifiable CSS or JS file and rewrite the URL for the minify handler.

Performance Enhancements

Minification

Static .js and .css requests coming from the CDN can optionally be minified before being delivered. This will result in decrease overall page size and improve load times.

This feature can be enabled/disabled in the /app_config/include/SitecoreCDN.config

```
<cdn enabled="true" ... minifyEnabled="true" ...>
```

CSS Processing

If minification is enabled, the media handler will also do URL replacements on css url([url]) statements. This ensures that any media referenced within the css also gets the benefit of the rewritten media urls.

This feature can be enabled/disabled in the /app_config/include/ SitecoreCDN.config

```
<cdn enabled="true" ... processCss="true" ...>
```

JS Fastloading

When an .aspx page request is having its urls rewritten, you have the option to automatically move any <script> tags to just before the </body> tag. This allows the browser to parallel load scripts rather than the browser imposed limit of 2 simultaneous downloads of script tags located in the <head>.

This feature can be enabled/disabled in the /app_config/include/ SitecoreCDN.config

```
<cdn enabled="true" ... fastLoadJsEnabled="true" ...>
```

Configuration

Under /App_Config/Include/ SitecoreCDN.config

```
<cdn enabled="true" filenameVersioningEnabled="true" minifyEnabled="true"
fastLoadJsEnabled="true" processCss="true">
  <provider type=" NTTData.SitecoreCDN.Providers.CDNProvider, NTTData.SitecoreCDN" />
  <!-- Incoming requests matching these urls will be processed
        .aspx is processed when Sitecore.Context.Item is resolved
  -->
  <processRequests>
    <regex pattern = "\.asmx" />
    <!-- matches any .asmx -->
  </processRequests>

  <!-- Incoming requests matching these urls will not be processed -->
  <excludeRequests>
    <!--<regex pattern = "Default\.aspx" />-->
  </excludeRequests>

  <!-- These regex patterns will prevent matching urls from being replaced in the outgoing html,
doesn't affect Incoming request processing -->
  <excludeUrls>
    <regex pattern = "\.axd" />
    <!-- this keeps ScriptResource.axd and WebResource.axd from being CDN'd-->
    <regex pattern = "VisitorIdentification.aspx" />
    <!-- this keeps the Sitecore Analytics request from being CDN'd -->
  </excludeUrls>
</cdn>
```

On the <cdn> element, `enabled="true"` turns the HTML rewriting on and off. `filenameVersioningEnabled="true"` turns on/off the versioning features of the URL rewrite.

`minifyEnabled = "true"` turns on/off the css/js minification.

`fastLoadJsEnabled = "true"` turns on/off the javascript tag relocation.

`processCss = "true"` turns on/off url rewriting within css files

You can override any behavior of the CDNManager class by implementing your own CDNProvider and changing `<provider type="NTTData.SitecoreCDN.Providers.CDNProvider, NTTData.SitecoreCDN" />`

You can specify which incoming requests will get Url replacement in <processRequests> by matching regex. By default, “.aspx” is only processed when Sitecore.Context.Item is not null.

You can specify which incoming requests to exclude from Url replacement in <excludeRequests> by matching regex.

Even when a request is being processed, you can specify a list of URLs that will not be replaced under <excludeUrls> by matching regex. Some media item you want going directly to the origin server. For example: Robot identifying files for DMS and .net generated script .axd files.